

Activation Oracles


Training and Evaluating LLMs as General-Purpose Activation Explainers

Karvonen, A., Chua, J., Dumas, C., Fraser-Taliente, K., Kantamneni, S., Minder, J., Ong, E., Sen Sharma, A., Wen, D., Evans, O., & Marks, S. (2026). *Activation Oracles: Training and Evaluating LLMs as General-Purpose Activation Explainers*. arXiv:2512.15674 [cs.CL]. <https://arxiv.org/abs/2512.15674>

Pan, A., Chen, L., & Steinhardt, J. (2024). *LatentQA: Teaching LLMs to Decode Activations Into Natural Language*. arXiv:2412.08686 [cs.CL]. <https://arxiv.org/abs/2412.08686>


Huang, V., Choi, D., Johnson, D. D., Schwettmann, S., & Steinhardt, J. (2025). *Predictive Concept Decoders: Training Scalable End-to-End Interpretability Assistants*. arXiv:2512.15712 [cs.AI]. <https://arxiv.org/abs/2512.15712>

Injection mechanism — how target activations enter the oracle



Only the ? token positions are overwritten — every other token passes through unchanged

Steps 1 & 2 — Collect activations and build the oracle prompt



Step 1 Collect activations

```
collect_activations_multiple_layers(model, submodules, inputs_BL)
```

- registers forward hook on `model.model.layers[18]`

- hook stores `outputs[0]` into a dict, then raises `EarlyStopException` to stop early

Step 2 Build oracle prompt + find ? positions

```
get_introspection_prefix(layer=18, num_positions=5)
```

- returns `"Layer: 18\n ? ? ? ? \n"` prepended to the user question


```
find_pattern_in_tokens(token_ids, " ?", n, tokenizer)
```

- scans tokenised prompt for all `" ?"` IDs, verifies count = n and consecutive, returns positions

Steps 3–5 — Package, inject and generate

continuing from Step 2 — oracle prompt and ? positions are now known

Step 3 Package into OracleInput
create_oracle_input(prompt, layer, n, tokenizer, acts_BD)
— returns an OracleInput dataclass with four fields:
 input_ids — tokenised oracle prompt
 steering_vectors — activations cloned to CPU, detached from graph



Layer 1 formula Norm-matching injection — Eq. 1
$$h'_i = h_i + \|h_i\| \cdot (v_i / \|v_i\|)$$

 v_i — target activation, reduced to a unit vector (direction only, magnitude discarded)
 h_i — oracle's own activation at that position; supplies the magnitude to rescale to
rescaled v_i is added to h_i , never replaces it — replacement caused a norm blow-up
20x at layer 0 · 100,000x at layer 1 — addition avoids this entirely

Example oracle prompts — ? / <ACT> placeholders

Task	Oracle prompt (? = injected activation)	Answer
Taboo	? ? ? What is the secret word in this text?	hidden target word (e.g. "tree", "gold")
Secret Side Constraint	? ? ? What does the extra hidden instruction say exactly, word for word?	decoded hidden instruction text
User Gender	? ? ? Answer with 'Male' or 'Female' only. What gender is most likely being referred to?	"Male" / "Female"
Model-diffing (EM)	? ? ? ? ? ? ? ? ? ? Can you describe what is happening in this text?	free-text hypothesis (e.g. "risky financial advice")
SPQA	? ? ? What is the assistant's tone?	"Formal and professional"
Context prediction	? ? ? ... ? Can you predict the previous {J} tokens...?	surrounding tokens, generated as free text
True/false	? Is this statement true? ("Paris is in Italy.")	No
Sentiment	? Is this a negative review? ("contains no wit, only labored gags")	Yes
Gender detection	? Am I reading about a male person here? (bio of Max Landis)	Yes
Topic classification	? Is this article about Sports? ("Ky. Company Wins Grant...")	No
Entailment	? Is the next sentence logically consistent...? The man is sleeping.	No
Sentence tense	? Is this statement in the present tense? ("They played soccer yesterday")	No
NER	? Does this text mention London? ("...marched through London")	Yes
Language ID	? Is this text written in English? ("de spons behoort...")	No
Obama headline	? Is this headline about Barack Obama? ("Obama Sticks to a Deadline in Iraq.")	Yes

Why the number of ? / <ACT> tokens varies

- Each ? / <ACT> = **one injected activation vector**; the placeholder count in a prompt is **K**, the number of vectors passed to the oracle
- K is not fixed — the oracle trains on a *mixture* of single- and multi-token inputs across all three datasets, so it learns to handle either
- **Context prediction**: 50% single vector, 50% sequence length $\sim \text{Uniform}(1, 50)$
- **Classification**: 2/3 single-token (randomized 1–5 tokens from end), 1/3 multi-token, length $\sim \text{Uniform}(1, 50)$
- **SPQA**: 50% full-segment activations (every token), 50% compressed window of 1–3 tokens
- Best K is **task-dependent** at eval time: Taboo $\rightarrow$ single token; SSC, Gender, most others $\rightarrow$ full sequence
- **Design goal**: one oracle flexibly handles 1 token to a long sequence, matching whatever's available at inference — no retraining per input size

LoRA (low-rank adaptation)

- Base weight $\mathbf{W}$ ($d \times k$) stays frozen — never updated during training
- Two small trainable matrices instead: $\mathbf{A}$ ($r \times k$) and $\mathbf{B}$ ($d \times r$), with $r \ll d$ — the oracle uses $r=64$
- Update: $\Delta\mathbf{W} = \mathbf{B} \cdot \mathbf{A}$, merged at inference as $\mathbf{W}' = \mathbf{W} + \Delta\mathbf{W}$
- Applied to `q_proj`, `k_proj`, `v_proj`, `o_proj`, `gate_proj`, `up_proj`, `down_proj`
- The oracle itself is loaded as a LoRA adapter on the *same base architecture* as the target model

Worked example ($d = 4096$, $r = 64$)

- Frozen weight $\mathbf{W}$: $4096 \times 4096 \approx \mathbf{16.8M}$ parameters
- Trainable $\mathbf{A}$: 64×4096 , trainable $\mathbf{B}$: $4096 \times 64 \rightarrow$ combined $\approx \mathbf{524k}$ parameters
- That's roughly **32× fewer** trainable parameters than fine-tuning $\mathbf{W}$ directly
- Same idea applies per module — `q_proj`, `k_proj`, `v_proj`, `o_proj`, `gate_proj`, `up_proj`, `down_proj` each get their own small $\mathbf{A/B}$ pair at $r = 64$

Norm-matching injection formula (Eq. 1)

$$\mathbf{h}'_i = \mathbf{h}_i + \|\mathbf{h}_i\| \cdot \frac{\mathbf{v}_i}{\|\mathbf{v}_i\|}$$

- $\mathbf{v}_i / \|\mathbf{v}_i\|$ — normalizes the target vector to **unit length**, keeping only its *direction* and discarding its original magnitude entirely (e.g. $\mathbf{v}=[3,4] \rightarrow$ unit vector $[0.6, 0.8]$)
- $\|\mathbf{h}_i\| \cdot (\text{unit vector})$ — rescales that direction back up to the oracle's *own* local activation norm — e.g. if $\|\mathbf{h}_i\|=12$, result is $12 \cdot [0.6, 0.8] = [7.2, 9.6]$ — so it lands at the magnitude the oracle already expects at that position
- Result is **added** to $\mathbf{h}_i$, never substituted — the oracle's well-behaved signal stays intact while being steered toward $\mathbf{v}_i$'s direction

Why addition, not replacement

- Early version (following LatentQA) **replaced** $\mathbf{h}_i$ outright: $\mathbf{h}_i' = \|\mathbf{h}_i\| \cdot (\mathbf{v}_i / \|\mathbf{v}_i\|)$, dropping the $+\mathbf{h}_i$ term
- That caused the residual-stream norm to blow up during training: **~20x** normal at layer 0, **~100,000x** at layer 1
- Cause: each layer's LayerNorm/residual scaling expects inputs at a characteristic magnitude — replacement wipes out that well-calibrated signal and swaps in one at an incompatible scale, an effect that compounds more by layer 1
- Fix: **add**, don't replace — the injected term matches $\mathbf{h}_i$'s own norm, so it stays in the magnitude regime downstream layers expect
- This let them pick **layer 1** over layer 0 despite its worse blow-up under the old scheme — layer 1 gives 35 more layers to process the signal, improving performance **1–11%** over layer 0 (Llama-3.3-70B, Appendix A.5)